

Project Report — Original Language Detection

*Students: M. Denend, J. Hoffmann, L. Prakash**Prof. Ray Mooney***Abstract**

Abstract goes here.

1 Introduction

Our task was to determine the original language of translated text. Determining the language of origin for written text is something that may potentially have applications for national security, for example, if terrorists use social media. Also, it may be of great importance to historians in order to trace back the path in which ancient novels may have been translated from, for example, the Bible. Doing so and using improving machine translation to translate a book back into prior languages may possibly open up avenues of research to better understand intended meanings of the first authors, as meaning can get lost when novels are translated. Furthermore, previous studies have shown that machine translators are much more accurate in identifying whether writing was from an original language, or translated, as mentioned in (I KNOW WHICH PAPER, I JUST NEED TO ACTUALLY CITE IT). Our goal was to build upon past research. In particular, we found that [1] had some good results when trying to determine the original language, however, it only tested against slices of the same books, which we felt defeats the potential for applying to other books. So we decided to have our train set and test set be from completely separate novels, which is one of our major contributions. We liked the idea of using horizontal slices of parse trees mentioned in [3], so we decided to take some features from [1] and add parse features to this. We built our dataset in a manner like in [1], but decided to extend it by adding a couple more books and two more languages. Our results were... (Jessica?)

2 Problem Definition

(a) Description

Our project is to take novels, freely available from Project Gutenberg (CITE THIS), and translated from a wide variety of languages, and determine which language they originated from. To do this, we take a novel electronically available in plain text format, and run some algorithms to determine various features that we detect from the books. We group these into three types of features (to be elaborated upon more later): lexical features, parse features (horizontal slices), and homemade features. Once we've computed these features for both our train and test sets, we feed them to a machine learning algorithm to see if it can determine the source language for the novels in the test set.

(b) Algorithm

We take a text format novel from Project Gutenberg. We clean out of it Chapter headings, subheaders following chapters, beginning and end informational text, and anything else that isn't words that we can use to compute. From here, we use the Stanford Parser's PCFG (CITE THIS) to build best-guess parse trees (in Penn Tree format) for the sentences in this novel, and we use the Stanford Parser's tokenizer to build tokenized sentences and detect ends of sentences. Both the tree format and text format are used in our machine learning, but we could just use the tree if space efficiency was a problem. The text format is used to

compute word unigrams (using Weka’s toolkit to do so), and the tree format is used to compute our parse features and homemade features as well. Both the parse and homemade feature computations were built from scratch with the ideas from [1] and [3], with some help from Stanford’s Parser package to read the Penn Trees. Finally, all the features were combined together and trained/tested upon using our (which?) machine learning algorithm.

(c) Methodology

From Project Gutenberg, we built a dataset of books written in different languages, and translated into English. We then cut these books into slices of 400 sentences, and compute features on these slices. After feature extraction, we feed these features into a classifier and test it on a completely different set of books.

3 Building the dataset

(a) Pitfalls

Since we built it ourselves, our dataset is rather small. We therefore were extremely careful not to overfit on books particularities that are not specific to the language, such as topic, style, author, geographic places, time period. We believe the way we’ve constructed the dataset satisfies these criteria.

(b) Description of the dataset

When building the dataset, we used books written from 1700 to 1920 to control the time period as best as we can. We made sure that we never used the same author nor translator (?) in train and test, and we tried varying the topics between train and test as much as we could once the previous criteria were met. When using Lexical Features, we only remembered the most common 1000 English words to control overfitting on topics as well. The dataset we used is further described on this table:

Language	Title	Author	Pub.Date
American	The Adventures of Huckleberry Finn	Mark Twain	1884
American	The Invisible Man	Orson Wells	1897
American	The Scarlet Letter	Nathaniel Hawthorne	1850
American	Moby Dick	Herman Melville	1851
American	Uncle Toms Cabin	Harriet Beecher Stowe	1852
American	The Narrative of A. G. Pym of Nantucket	Edgar Allen Poe	1899
English	Great Expectations	Charles Dickens	1861
English	The Picture of Dorian Gray	Oscar Wilde	1891
English	Jude the Obscure	Thomas Hardy	1895
English	Treasure Island	R.L Stevenson	1883
English	Middlemarch	George Eliot (M. Evans)	1874
English	Frankenstein	Mary Wollstonecraft Shelley	1818
French	Around the world in 80 days	Jules Verne	1873
French	Father Goriot	Balzac	1835
French	Les Miserables	Victor Hugo	1862
French	Madame Bovary	Gustave Flaubert	1857
French	The Man in the Iron Mask	Alexandre Dumas	1637
French	Candide	Voltaire	1759
French	The Red and the Black	Stendahl	1830
Russian	Fathers and Children	Ivan Turgenev	1862
Russian	The Man who was afraid	Maxim Gorky	1899
Russian	War and Peace	Leo Tolstoy	1869
Russian	The Brothers Karamazov	Fyodor Dostoyevsky	1880
Russian	A hero of Our Time	M. Y. Lermontov	1840
Russian	The Death of the Gods	Dmitri Mérejkowski	1895
Spanish	Don Quixote	Miguel de Cervantes	1605
Spanish	Dona Perfecta	Benito Pérez Galdós	1876
Spanish	The Visions of Quevedo	Francisco de Quevedo	1626
Spanish	Tragic Sense of Life	Miguel de Unamuno	1912
Spanish	The Life of Lazarillo of Tormes	Lazarillo of Tormes	1554
Spanish	The Four Horsemen of the Apocalypse	Vicente Blasco Ibanez	1916
Spanish	Pepita Ximenez	Juan Valera	1874
Spanish	The Fourth Estate	Armando Palacio Valdés	1901
German	The Devil's Elixir	E. T. A. Hoffmann	1815
German	Debit and Credit	Gustave Freytag	1855
German	Old Fritz and the New Era	Luise Mühlbach	1868
German	Venus in Furs	Leopold V. Sacher-Masoch	1870
German	The Sorrows of Young Werther	Johann Wolfgang von Goethe	1774
German	The Banished: A Swabian Historical Tale	Wilhelm Hauff	1839

4 Features

(a) General Description

When computing the features, we cut the books in slices of n sentences. With $n = 1$, we get a lot of training examples, but finding out the original language from only 1 sentence seems unreasonable (which was proven by experiment). With the whole book as a sample, we didn't get enough training instances. We chose $n = 400$, as it provided good balance between number of train samples, and informativeness of the samples.

In training, we take 8 slices of each book, so that we have the same size of training for each language. In testing, we take all the slices, except the incomplete last one.

(b) Lexical Feature

Three different kind of lexical features were used:

1000 most common English words After overfitting on words like "Moscow" or "Madame" when using unigrams, we decided to only use the 1000 most common words in English to avoid overfitting on topics, local idioms, or geographic places.

POS unigram We only used POS unigrams, since we believed using bigram or trigram would just add redundancy with the Parse features.

Etymology For each word used, we computed the etymology of the word from parsing Wiktionary [2], and added all possible etymology as a feature. Rational being that if you're translating from a latin language, you may be more likely to use a latin word when choosing between synonyms. We only included the etymologies of nouns, verb, adjective and adverb, since we believed preposition, pronouns and others wouldn't be significant and would only drown the relevant data.

(c) Parse Features

Inspired by [3], we ran the Stanford parser on our slices, and computed the parse trees. By running a BFS, we then compute the generation rules and count them. We don't take into account the leaves of the trees to avoid redundancy with lexical features (and possible topic overfitting).

(d) Homemade Features

We also computed all the features of [1], which include length of sentences, length of words, ratio of pronouns, nouns, complex sentences etc. [MATT ADD WHATEVER YOU WANT].

5 Classification task

(a) Features Extraction

(b) Comparisons of different algorithm

6 Experiments

(a) Description of different experiments

We conducted the following experiments:

- using unigrams (overfitting on topics)
- using the same books for train and test

On top of that, we also inquired the effect of different features on the dataset:

- Most common words only
- Most common + etymology
- POS only
- All lexical features (Most common words + etymology + POS)
- Homemade only
- Parse only
- Most common + etymology + Homemade
- All features

(b) Results

Experiment	Accuracy	Best classifier
1000 Most common words	96.67 %	LR

Figure 1: Testing and training on the same books (different parts)

As expected, using the same books for training and testing gives excellent results even when using only the 1000 most common English words. Since we used a lot of books in common with [1], we’re wondering how they got such poor results.

Experiment	Accuracy
Unigrams	61.80 %
1000 Most common words	62.92 %
1000 Most common words + etymology	61.80 %
1000 Most common words + POS	61.80 %
Parse features only	

Figure 2: Lexical Features

We get better results when using only the 1000 most common English words, as opposed to all the words (Unigrams). This was a very positive result, since it suggests that we did build our dataset in a robust way. Indeed, using unigrams completely overfitted on topic, as illustrated by the following example: we used Moby Dick in our training, and therefore *whale* is the third most relevant feature for American books. Contrary to intuition, the etymology feature didn’t help at all.

7 Conclusion

References

- [1] Gerard Lynch and Carl Vogel. Towards the automatic detection of the source language of a literary translation. 2012.
- [2] Christian M. Meyer and Iryna Gurevych. Wiktionary: A new rival for expert-built lexicons? Exploring the possibilities of collaborative lexicography. In Sylviane Granger and Magali Paquot, editors, *Electronic Lexicography*, chapter 13, pages 259–291. Oxford: Oxford University Press, November 2012.

- [3] Sze-Meng Jojo Wong and Mark Dras. Exploiting parse structures for native language identification. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pages 1600–1610. Association for Computational Linguistics, 2011.